

PartA

Problem1:

Q1:

(habitat) Singularity> ./vecadd00 500

Total vector size: 3840000

Time: 0.000331 (sec), GFlopsS: 11.603838, GBytesS: 139.246058

Test PASSED

(habitat) Singularity> ./vecadd00 1000

Total vector size: 7680000

Time: 0.000589 (sec), GFlopsS: 13.041399, GBytesS: 156.496784

Test PASSED

(habitat) Singularity> ./vecadd00 2000

Total vector size: 15360000

Time: 0.001138 (sec), GFlopsS: 13.497697, GBytesS: 161.972368

Test PASSED

Q2:

(habitat) Singularity> ./vecadd01 500

Total vector size: 3840000

Time: 0.000008 (sec), GFlopsS: 488.064465, GBytesS: 5856.773585

Test FAILED

(habitat) Singularity> ./vecadd01 1000

Total vector size: 7680000

Time: 0.000008 (sec), GFlopsS: 947.419256, GBytesS: 11369.031078

Test FAILED

(habitat) Singularity> ./vecadd01 2000

Total vector size: 15360000

Time: 0.000008 (sec), GFlopsS: 1894.838513, GBytesS: 22738.062155

Test FAILED

It appears that `vecadd01` consistently fails with shorter execution times but significantly larger values for GFlopsS and GBytesS.

Problem2:

Q3:

(habitat) Singularity> ./matmult00 256

Data dimensions: 4096x4096

Grid Dimensions: 256x256

Block Dimensions: 16x16

Footprint Dimensions: 16x16

Time: 0.089620 (sec), nFlops: 137438953472, GFlopsS: 1533.572636

(habitat) Singularity> ./matmult00 512

Data dimensions: 8192x8192

Grid Dimensions: 512x512

Block Dimensions: 16x16

Footprint Dimensions: 16x16

Time: 0.690086 (sec), nFlops: 1099511627776, GFlopsS: 1593.296787

TEST FAILED: number of errors: 4321886, max rel error: 0.000122

(habitat) Singularity> ./matmult00 1024

Data dimensions: 16384x16384

Grid Dimensions: 1024x1024

Block Dimensions: 16x16

Footprint Dimensions: 16x16

Time: 5.569285 (sec), nFlops: 8796093022208, GFlopsS: 1579.393612

TEST FAILED: number of errors: 138340737, max rel error: 0.000244

(habitat) Singularity>

Q4:

(habitat) Singularity> ./matmult00 256

Data dimensions: 8192x8192

Grid Dimensions: 256x256

Block Dimensions: 32x32

Footprint Dimensions: 32x32

Time: 0.554000 (sec), nFlops: 1099511627776, GFlopsS: 1984.678384

TEST FAILED: number of errors: 4321886, max rel error: 0.000122

(habitat) Singularity> ./matmult00 512

Data dimensions: 16384x16384

Grid Dimensions: 512x512

Block Dimensions: 32x32

Footprint Dimensions: 32x32

Time: 4.439757 (sec), nFlops: 8796093022208, GFlopsS: 1981.210521

TEST FAILED: number of errors: 138340737, max rel error: 0.000244

(habitat) Singularity> ./matmult00 1024

Data dimensions: 32768x32768

Grid Dimensions: 1024x1024

Block Dimensions: 32x32

Footprint Dimensions: 32x32

Time: 35.582783 (sec), nFlops: 70368744177664, GFlopsS: 1977.606522

TEST FAILED: number of errors: 849488718, max rel error: 0.000488
(habitat) Singularity>

Q5:

The experiments show that optimizing memory access patterns and choosing appropriate thread block sizes significantly improve CUDA program performance, while highlighting the importance of numerical stability and correctness considerations.

Part-B: CUDA Unified Memory (20 points)

In this problem we will compare vector operations executed on host vs on GPU to quantify the speed-up.

Q1 4 points Write a C++ program that adds the elements of two arrays with a K million elements each. Here K is a parameter. Profile and get the time to execute this program for K=1,5,10,50,100. Use free() to free the memory at the end of the program.

(habitat) Singularity> g++ -o Q1 Q1.cpp

(habitat) Singularity> ./Q1

Time to execute for K = 1 million: 4 ms

Time to execute for K = 5 million: 20 ms

Time to execute for K = 10 million: 41 ms

Time to execute for K = 50 million: 205 ms

Time to execute for K = 100 million: 409 ms

Q2:

(habitat) Singularity> ./Q2

K = 1, 1 Block, 1 Thread, Time: 76 ms

K = 1, 1 Block, 256 Threads, Time: 9 ms

K = 1, 3907 Blocks, 256 Threads/Block, Time: 22 ms

K = 5, 1 Block, 1 Thread, Time: 90 ms

K = 5, 1 Block, 256 Threads, Time: 8 ms

K = 5, 19532 Blocks, 256 Threads/Block, Time: 85 ms

K = 10, 1 Block, 1 Thread, Time: 77 ms

K = 10, 1 Block, 256 Threads, Time: 7 ms

K = 10, 39063 Blocks, 256 Threads/Block, Time: 158 ms

K = 50, 1 Block, 1 Thread, Time: 76 ms

K = 50, 1 Block, 256 Threads, Time: 8 ms

K = 50, 195313 Blocks, 256 Threads/Block, Time: 760 ms

K = 100, 1 Block, 1 Thread, Time: 49 ms

K = 100, 1 Block, 256 Threads, Time: 7 ms

K = 100, 390625 Blocks, 256 Threads/Block, Time: 1501 ms

(habitat) Singularity>

Q3:

(habitat) Singularity> ./Q3

Time to execute for K = 1 million: 42 ms

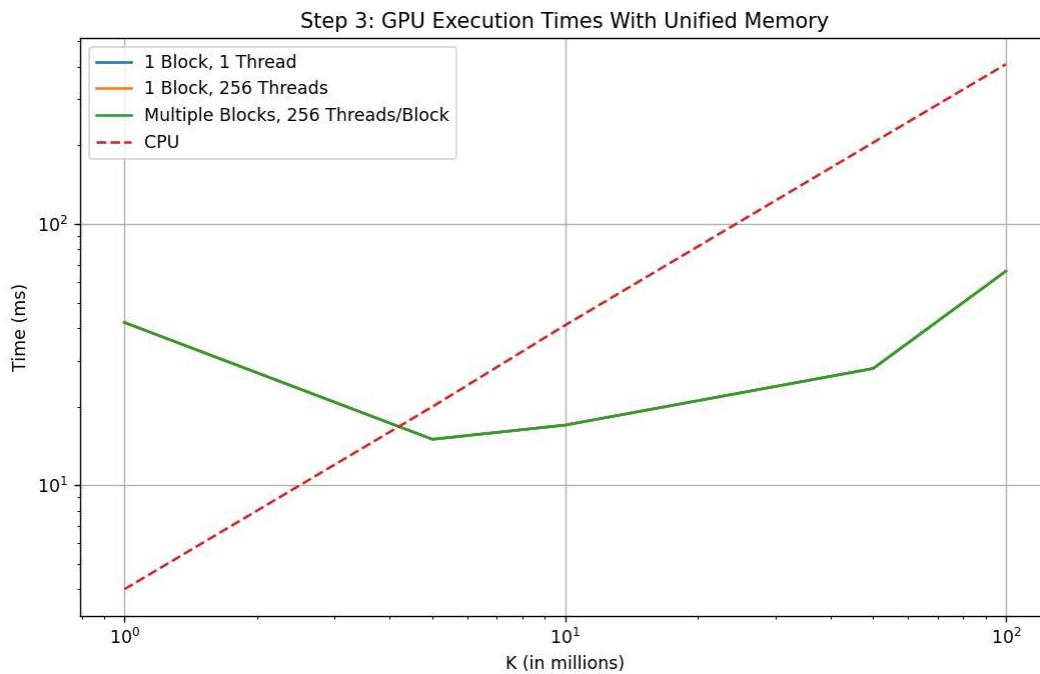

Time to execute for K = 5 million: 15 ms

Time to execute for K = 10 million: 17 ms

Time to execute for K = 50 million: 28 ms

Time to execute for K = 100 million: 66 ms

Q4:



Part-C: Convolution in CUDA (40 points)

C1:

(habitat) Singularity> ./c1

Kernel execution time: 11.012 milliseconds

Checksum: 122756344698240.000

C2:

(habitat) Singularity> ./c2

Tiled convolution kernel execution time: 1.925 milliseconds

Checksum of the tiled output tensor O: 122756344698240.000

C3:

(habitat) Singularity> ./c3

Iteration 0: Convolution execution time: 8015.012 milliseconds

Iteration 1: Convolution execution time: 13.058 milliseconds

Iteration 2: Convolution execution time: 13.016 milliseconds

Iteration 3: Convolution execution time: 12.997 milliseconds

Iteration 4: Convolution execution time: 12.993 milliseconds

Checksum: 122756344698240.000